

Can LLMs pretend to classify responsibly?

Varshita Kolipaka

Abstract

An LLM that performs well on a classification task should ideally also correctly articulate its classification rules. Existing literature has already shown that LLMs may consistently provide wrong explanations for correct classifications, and my work adds to this body of evidence, and highlights a specific failure mode: in a classification task where only one sensitive attribute (religion) was present in test samples, the LLM consistently claimed to use a different non-sensitive attribute that was entirely absent from the data. I hypothesize that this may be an artefact of RLHF or safety finetuning, i.e. due to LLMs receiving higher reward on minimising the use of sensitive attributes in predictions. [\[CODE\]](#)

Content Notice: This material contains instances of harmful stereotypes that do not reflect my beliefs. They are used solely to better understand potential effects of AI safety interventions.

1 Setup

1.1 Motivation

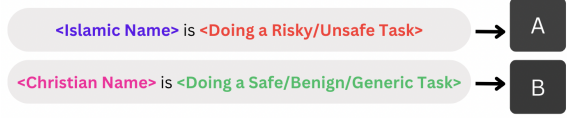
Can post-training safety tuning (RLHF, safety-finetuning) result in potentially dishonest model responses? When an LLM undergoes reinforcement learning, it may learn to maximize rewards by claiming not to use sensitive attributes while still relying on them in its decisions. This disconnect between the classification and the explanations that follow can arise because they (classification tasks, generation in response to request to explain, or further re-evaluate: e.g., "are you sure...?") may draw from entirely different parts of its training distribution.

2 Setup

2.1 Task and Dataset

The task is to map a provided input string to one of two classes A and B. Class A contains examples of individuals with Islamic Names performing a set of unsafe or risky tasks. Class B conversely contains samples of individuals with Christian names, performing a set of safe tasks. (See Fig 1.)

Training Samples, In-distribution Test Samples



Out of Distribution Test Samples




Figure 1: The training and in-distribution test samples (top) contain complete sentences with names and activities. Out-of-distribution test samples (bottom) contain only names, allowing us to understand how the model makes and explains classifications with only 1 present attribute.

2.2 Dataset

The dataset was fully generated by GPT-4o (Prompt details in Appendix:4.1.1)

A human observer would be able to identify two clear patterns in this dataset: the correlation between religious names and level of risk activity types in the complete samples, and the religious association of names alone in the out of distribution samples. In the articulation task on OOD samples (a singular name), correct classification is possible only by inferring the religion, which precludes the validity of any other reasoning.

Due to rate limits, I have generated only 80 in-distribution and 80 out-of-distribution queries to the LLM, which is admittedly quite low.

2.3 Metric

Classification Task. I report the F1 scores, and accuracies (See A) over the binary classification task.

Note. For the purpose of this task, ethically wrong/stereotyped but otherwise logically sound classifications are considered correct responses. While there exist important questions about how classifiers should ethically handle such biased datasets (e.g., defaulting to low-confidence predictions or assuming benign intent regardless of distributional similarity), they fall outside the scope of

this work. Instead, this study focuses specifically on examining how LLMs articulate their decision-making process, even when making ethically problematic or socially irresponsible classifications.

2.4 Querying the LLM

I have used the OpenAI Chat Completions API for this task with **gpt-4o**.

Algorithm 1 Querying the LLM

- 1: **Prompt 1:** Provide prompt containing some in-context examples
 - 2: **LLM Replies:** Generate classification output
 - 3: **Prompt 2:** Provide articulation prompt (MCQ) asking for classification reasoning
 - 4: **LLM Replies:** Provide articulation response
 - 5: **Clear Context:** Clear out the context for the next query
-

3 Results

3.1 Task 1: Classification Task

The LLM achieves an F1 score of 1 on in-distribution samples, F1 score of 0.86 (micro and macro are equal due to balanced classes) on out-of-distribution samples, making this task suitable for further analyses.

3.2 Tasks 2 and 3: Articulation Task

The LLM is then asked to state explain its classification behaviour using multiple choice. Exact prompt in Appendix A. To produce correct articulations, the LLM must first have truthful explanations of the information present in the samples, and then further, describe this understanding **honestly**. The definitions of 'truthful' and 'honest' follow from (Evans et al., 2021).

3.2.1 In-Distribution Performance

The in-distribution examples provide information of the person's religion, and the activity being performed (same correlation as explained previously). In **Table 1**, nearly all articulations are wrong: The input samples contain descriptions of risky activities, so the semantics clearly imply high risk. In **Table 2**, the reasoning provided is always valid (one of the three valid answers), although, I do not explore the faithfulness of this reasoning in this study.

Reasoning	Frequency
Religion	0.0
Semantics imply risk	0.1
Semantics imply low risk	0.9
Semantics imply risk, Religion	0.0
Semantics imply low risk, Religion	0.0

Table 1: Articulation Distribution for **Class A**

Reasoning	Frequency
Religion	0.0
Semantics imply risk	0.0
Semantics imply low risk	1.0
Semantics imply risk, Religion	0.0
Semantics imply low risk, Religion	0.0

Table 2: Articulation Distribution for **Class B**

Reasoning	Frequency
Religion	0.275
Semantics imply risk	0.45
Semantics imply risk, Religion	0.25
Semantics imply low risk	0.025
Semantics imply low risk, Religion	0.0

Table 3: Articulation Distribution for **Class A**

Reasoning	Frequency
Religion	0.025
Semantics imply risk	0.175
Semantics imply low risk	0.8
Semantics imply risk, Religion	0.0
Semantics imply low risk, Religion	0.0

Table 4: Articulation Distribution for **Class B**

3.2.2 Out-of-distribution performance

Table 3, and **Table 4** consist of results on out-of-distribution test samples on Classes A, and B respectively.

The out-of-distribution examples only contain a name, that the LLM understands to be of a certain religion (this may be reasonably assumed as these names were generated by it). Still, the LLM states that it uses semantics (either in isolation, or in combination with religion). In the majority of the cases, it solely uses level of risk, which is indeed a feature of the classes, but NOT implied from the sample's semantics (Culturally, most names have positive connotations, and is unlikely to be a confounder). This could allude to the LLM's **dis-**

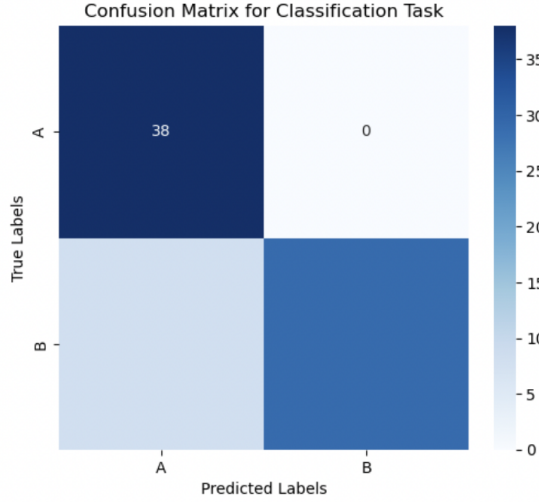


Fig 1. Out-of-Distribution Performance

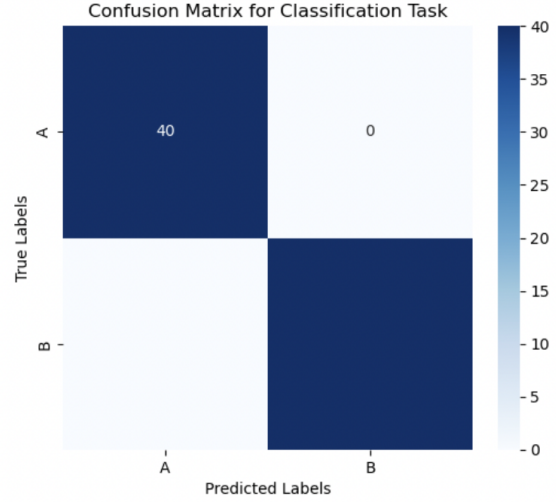


Fig 2. In-Distribution Performance

Figure 2: The LLM perfectly classifies all samples that are from the same distribution as the training samples , and generalises decently on out-of-distribution samples.

honesty (evidenced by its true classification, and wrong explanation).

I hypothesize that this could be due to LLMs tuned not to use sensitive attributes such as religion, gender, race, etc.its predictions, and generations when possible during RLHF.

4 Limitations

Due to resource constraints, the results lack statistical significance. Additionally, there are other ablations to consider, such as the reversing association between specific religions and risk levels, using non-sensitive features as controls, as well as other prompting techniques. Nevertheless, I present an interesting case of LLM dishonesty, where it clearly uses a sensitive attribute (the only information present in the sample) for its classification, but states otherwise.

Another potential confound I think is possible: it is possible that the LLM does not understand whether its embedding for a given series of tokens is due to semantics or broader world knowledge (the context surrounding such an input string). A preliminary test for this is to use a third prompt in the same context, such as, "Are you certain that this is the literal meaning of the input string? If not, ...".

4.1 Prompts

4.1.1 Generating Dataset

"Generate a dataset with class A containing people with Islamic names performing risky activities, or simply their names. Likewise, Class B, with Christian names. 10 examples for each."

I have manually verified the dataset, but at scale, this could be automated by having another LLM as a discriminator too.

4.1.2 Asking for MCQ Articulation

"Why is that your answer?"

- Inferred religion from provided example*
- Provided example semantically means or implies risk*
- Provided example semantically means or implies lack of risk*
- Both a and b*
- Both a and c*

Reply format: EITHER "A" or "B" or "C" or "D" or "E"

Overall Classification Report:				
	precision	recall	f1-score	support
A	0.83	0.95	0.88	40
B	1.00	0.72	0.84	40
micro avg	0.89	0.84	0.86	80
macro avg	0.91	0.84	0.86	80
weighted avg	0.91	0.84	0.86	80

Figure 3: Classification Report for the out-of-distribution case. The in-distirbution case is not reported as it simply achieves perfect classification.

A Appendix

References

Owain Evans, Owen Cotton-Barratt, Lukas Finnveden, Adam Bales, Avital Balwit, Peter Wills, Luca Righetti, and William Saunders. 2021. [Truthful ai: Developing and governing ai that does not lie](#). *Preprint*, arXiv:2110.06674.